


Learning

Learning

- The learning is an important feature of human **computational ability**.
- Learning may be viewed as the **change in behavior** acquired due to **practice or experience**, and it lasts for relatively long time.
- As it occurs, the **effective coupling** between the neuron is modified.
- In case of artificial neural networks, it is a **process of modifying neural network by updating its weights, biases and other parameters**, if any.
- During the learning, **the parameters of the networks are optimized and as a result process of curve fitting**. It is then said that the network has passed through a learning phase.
- Learning refers to the process by which an **algorithm** or **model improves** its performance on a task by:
 - **Analyzing data or**
 - **Interacting with an environment,**
 without being explicitly programmed for every scenario.

Machine Learning:

- **Arthur Samuel's Definition (1959)**: "Machine learning is the field of study that gives computers the ability to learn without being explicitly programmed".
- **Tom M. Mitchell's Definition (1997)**: "A computer program is said to learn from experience **E** with respect to some class of tasks **T** and performance measure **P**, if its performance at tasks in **T**, as measured by **P**, improves with experience **E**".



Classes of Learning Problems

Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn function to map
 $x \rightarrow y$

Apple example:



This thing is an apple.

Unsupervised Learning

Data: x

x is data, no labels!

Goal: Learn underlying structure

Apple example:



This thing is like the other thing.

Reinforcement Learning

Data: state-action pairs


Goal: Maximize future rewards over many time steps

Apple example:



Eat this thing because it will keep you alive.

Types of Machine Learning



Dataset

- A dataset is a **collection of related data** organized in a structured way, usually in the form of rows and columns, that is used for analysis, training machine learning models, or drawing conclusions.

Student ID	Name	Grade	Attendance (%)
101	Ayesha	85	92
102	Ravi	78	88
103	Fatima	91	95

- This table is a **dataset** where:
 - Each **row** = a data entry or record (101, Ayesha, 85, 92)
 - Each **column** = a feature (like Student ID, Name, Grade, Attendance (%))

Labelled Dataset

- A **labeled dataset** includes **input data** *and* the **correct output (label or target value)** for each example.

Square Feet	Bedrooms	Bathrooms	Location	Age (Years)	Price (\$)
1400	3	2	New York	10	550,000
1600	4	2	Chicago	5	460,000
1200	2	1	Dallas	15	300,000
1800	3	2	San Diego	7	620,000
2000	4	3	Seattle	2	750,000

- The first **5 columns** are the **features (inputs)**.
- The last column (**Price**) is the label (**output**) — this is what the model learns to predict.

Unlabelled Dataset

- An **unlabeled dataset** includes only **input data**, with **no labels** or correct answers provided.

Square Feet	Bedrooms	Bathrooms	Location	Age (Years)
1400	3	2	New York	10
1600	4	2	Chicago	5
1200	2	1	Dallas	15
1800	3	2	San Diego	7
2000	4	3	Seattle	2

- This dataset has **no labels (i.e., no price column)**.
- It includes **only input features**.

Independent Variable

- Also called as **input, feature, or predictors**.
- The variable that you **control** or use to **predict** something else.
- These are the columns that **affect or influence** the outcome.

*"What you **change or know** to try to predict something."*

- In house price prediction example, the **Independent Variables** are **Square Feet, Bedrooms, Location, Age**.

Dependent Variable

- Also called **label, target, or output**.
- The variable you are trying to **predict or explain**.
- It **depends on** the values of the independent variables.



*"What you want to **find out or predict**."*

- In house price prediction example, **Dependent Variable is Price**.
- Because the price **depends on** the size, number of bedrooms, location, and age.


Supervised Learning

- Supervised learning is a foundational methodology within the field of machine learning that focuses on learning a **mapping from input variables to an output variable** based on example **input-output pairs**.
- This approach relies on a **labelled dataset**, where each data point is associated with a corresponding **target value or label**.
- In this learning, **every input pattern** that is used to train the network is **associated with an output pattern** and that data is called training data.
- The primary objective is to **infer a function or model that can generalize from the training data to accurately predict outputs for unseen inputs**.
- Here, the output of a network is compared with the corresponding target value and the **error is determined**.

Supervised Learning



Key Characteristics of Supervised Learning



Types of Supervised Learning

Now, Supervised learning can be applied to two main types of problems:

- **Classification:** Where the output is a categorical variable.
 - Example: **spam vs. non-spam emails, yes vs. no.**
- **Regression:** Where the output is a continuous variable.
 - Example: **predicting house prices, stock prices.**

User ID	Gender	Age	Salary	Purchased	Temperature	Pressure	Relative Humidity	Wind Direction	Wind Speed
15624510	Male	19	19000	0	10.69261758	986.882019	54.19337313	195.7150879	3.278597116
15810944	Male	35	20000	1	13.59184184	987.8729248	48.0648859	189.2951202	2.909167767
15668575	Female	26	43000	0	17.70494885	988.1119385	39.11965597	192.9273834	2.973036289
15603246	Female	27	57000	0	20.95430404	987.8500366	30.66273218	202.0752869	2.965289593
15804002	Male	19	76000	1	22.9278274	987.2833862	26.06723423	210.6589203	2.798230886
15728773	Male	27	58000	1	24.04233986	986.2907104	23.46918024	221.1188507	2.627005816
15598044	Female	27	84000	0	24.41475295	985.2338867	22.25082295	233.7911987	2.448749781
15694829	Female	32	150000	1	23.93361956	984.8914795	22.35178837	244.3504333	2.454271793
15600575	Male	25	33000	1	22.68800023	984.8461304	23.7538641	253.0864716	2.418341875
15727311	Female	35	65000	0	20.56425726	984.8380737	27.07867944	264.5071106	2.318677425
15570769	Female	26	80000	1	17.76400389	985.4262085	33.54900114	280.7827454	2.343950987
15606274	Female	26	52000	0	11.25680746	988.9386597	53.74139903	68.15406036	1.650191426
15746139	Male	20	86000	1	14.37810685	989.6819458	40.70884681	72.62069702	1.553469896
15704987	Male	32	18000	0	18.45114201	990.2960205	30.85038484	71.70604706	1.005017161
15628972	Male	18	82000	0	22.54895853	989.9562988	22.81738811	44.66042709	0.264133632
15697686	Male	29	80000	0	24.23155922	988.796875	19.74790765	318.3214111	0.329656571
15733883	Male	47	25000	1					

Figure A: CLASSIFICATION

It is a dataset of a shopping store that is useful in predicting whether a customer will purchase a particular product under consideration or not based on his/ her gender, age, and salary.

Input: Gender, Age, Salary


Output: Purchased i.e. 0 or 1; 1 means yes the customer will purchase and 0 means that the customer won't purchase it.

Figure B: REGRESSION

It is a Meteorological dataset that serves the purpose of predicting wind speed based on different parameters.


Input: Dew Point, Temperature, Pressure, Relative Humidity, Wind Direction

Output: Wind Speed




Working of Supervised Learning

- Where **supervised learning algorithm** consists of input features and corresponding output labels. The process works through:
 - Training Data:** The model is provided with a training dataset that includes input data (features) and corresponding output data (labels or target variables).
 - Learning Process:** The algorithm processes the training data, learning the relationships between the input features and the output labels. This is achieved by adjusting the model's parameters to minimize the difference between its predictions and the actual labels.



- Training phase** involves feeding the algorithm **labelled data**, where each data point is paired with its correct output.
 - The algorithm learns to identify patterns and relationships between the input and output data.
- Testing phase** involves feeding the algorithm new, unseen data and evaluating its ability to predict the correct output based on the learned patterns.





Training a Supervised Learning Model


Training a model for supervised learning involves several crucial steps, each designed to prepare the model to make accurate predictions or decisions based on labeled data.

- 1. Data Collection and Preprocessing:** Gather a labeled dataset consisting of input features and target output labels. Clean the data, handle missing values, and scale features as needed to ensure high quality for **supervised learning** algorithms.
- 2. Data Splitting:** Divide the data into **training set (80%) and the test set (20%)**.
- 3. Choosing the Model:** Select appropriate algorithms based on the problem type. This step is crucial for effective **supervised learning** in AI.
- 4. Training the Model:** Feed the model input data and output labels, allowing it to learn patterns by adjusting internal parameters.
- 5. Evaluating the Model:** Test the trained model on the unseen test set and assess its performance using various metrics.
- 6. Hyper-Parameter Tuning:** Adjust settings that control the training process (e.g., learning rate) using techniques like grid search and cross-validation.
- 7. Final Model Selection and Testing:** Retrain the model on the complete dataset using the best hyperparameters testing its performance on the test set to ensure readiness for deployment.
- 8. Model Deployment:** Deploy the validated model to make predictions on new, unseen data.

Supervised Learning Flow




Supervised Learning Algorithms and Their Applications



Type of Supervised Learning Algorithms

Algorithm	Regression, Classification	Purpose	Method	Use Cases
Linear Regression	Regression	Predict continuous output values	Linear equation minimizing sum of squares of residuals	Predicting continuous values
Logistic Regression	Classification	Predict binary output variable	Logistic function transforming linear relationship	Binary classification tasks
Decision Trees	Both	Model decisions and outcomes	Tree-like structure with decisions and outcomes	Classification and Regression tasks
Random Forests	Both	Improve classification and regression accuracy	Combining multiple decision trees	Reducing overfitting, improving prediction accuracy
SVM	Both	Create hyperplane for classification or predict continuous values	Maximizing margin between classes or predicting continuous values	Classification and Regression tasks
KNN	Both	Predict class or value based on k closest neighbors	Finding k closest neighbors and predicting based on majority or average	Classification and Regression tasks, sensitive to noisy data
Gradient Boosting	Both	Combine weak learners to create strong model	Iteratively correcting errors with new models	Classification and Regression tasks to improve prediction accuracy
Naive Bayes	Classification	Predict class based on feature independence assumption	Bayes' theorem with feature independence assumption	Text classification, spam filtering, sentiment analysis, medical

Advantages and Disadvantages of Supervised Learning



Advantages and Disadvantages of Supervised Learning


Advantages	Disadvantages
Accurately predicts patterns and enables data-driven decisions.	Risk of overfitting : models may capture noise, reducing performance on new data.
Uses labeled training data to effectively learn input-output relationships.	Requires extensive feature engineering , which is time-consuming and needs domain expertise.
Handles both classification and regression tasks.	Susceptible to bias in models if training data is biased, leading to unfair predictions.
Supports complex applications like image recognition and natural language processing (NLP) .	Heavy dependence on labeled data , which is costly and time-consuming to collect.
Well-defined evaluation metrics (accuracy, precision, recall, F1-score) for model assessment.	Performance strongly depends on data quality and representativeness .
Can build complex models for accurate predictions on unseen data.	Scalability issues if data labeling is not feasible for large datasets.

Challenges of Supervised Learning

1

Need for Labeled Data

Labeled data is crucial but impacts performance heavily.



3

Limited Adaptability

Limited adaptability needs few resources and affects performance minimally.



2

Overfitting


Overfitting requires high resources and significantly impacts performance.

4

Computational Costs

Computational costs are high but impact performance less directly.

Applications of Supervised Learning



MCQs

1. Which are essential elements of supervised learning?

- i. Input features
- ii. Output labels
- iii. Only unlabeled data
- iv. A mapping function (model)

2. Which of the following are examples of supervised learning tasks?

- i. Predicting house prices based on size and location
- ii. Spam email classification
- iii. Customer segmentation without labels
- iv. Detecting hidden groups in social networks

3. Which datasets can be directly used for supervised learning?

- i. Images of animals without labels
- ii. Housing data with prices
- iii. Medical images with disease labels
- iv. Web click logs without categories

4. Which are true about training and testing in supervised learning?

- i. Training data is used to teach the model
- ii. Testing data is used to evaluate performance
- iii. Both training and testing use the same dataset without splitting
- iv. No labeled data is required for testing

5. Which tasks are classification problems?

- i. Predicting whether a review is positive or negative
- ii. Diagnosing patients as diabetic or non-diabetic
- iii. Predicting a house price in dollars
- iv. Estimating temperature tomorrow in °C

6. Which tasks are regression problems?

- i. Classifying an image as cat or dog
- ii. Predicting rainfall in mm
- iii. Estimating exam marks out of 100
- iv. Detecting whether a transaction is fraudulent

7. Which algorithms belong to supervised learning?

- i. Logistic Regression
- ii. Decision Trees
- iii. K-Means Clustering
- iv. Support Vector Machines

8. Which challenges occur in supervised learning?

- i. Overfitting to training data
- ii. Class imbalance (e.g., 90% healthy, 10% sick)
- iii. Lack of labels is never an issue
- iv. Need for large labeled datasets

9. Which metrics are appropriate for classification problems?

- i. Accuracy
- ii. Precision
- iii. Recall
- iv. Mean Squared Error (MSE)

10. Which metrics are used in regression problems?

- i. F1-score
- ii. Confusion Matrix
- iii. Mean Squared Error (MSE)
- iv. R^2 (Coefficient of Determination)

Supervised Learning Problems

Student Exam Performance

Suppose we collected the following dataset of students:

Hours Studied (X1)	Attendance % (X2)	Marks (Y)
2	60	45
4	70	55
5	80	65
6	85	70
8	90	85
9	95	92

Build a **supervised learning model** that predicts **student exam marks (Y)** based on **hours studied (X1)** and **attendance % (X2)**.

Q1: Identify Key Elements

- What are the **features (inputs)** in this dataset?
- What is the **label (output)**?
- Is this problem **classification** or **regression**? Why?

Q2: Data Preparation

- If you split the dataset into 80% training and 20% testing, how many rows go into each?
- Why do we need to split the data?

Q3: Algorithm Choice

- Which supervised learning algorithms could be applied here?
 - Linear Regression
 - Decision Trees
 - Logistic Regression
 - Support Vector Machine (SVM)
 - Which one makes the most sense, and why?

Q4: Prediction Task

- If a student studies for **7 hours** and has **88% attendance**, what would your model likely predict?
- How would the prediction differ if the student had only **50% attendance**?

Q5: Evaluation

- Which metrics are suitable for evaluating this model: Accuracy, Precision, Recall, Mean Squared Error (MSE), R^2 ?
- Why?

Q6: Challenges

- What problems might occur if the dataset was too small?
- What if we only collected “Hours Studied” but not “Attendance”?

Q7: Real-World Application

- How could schools or coaching institutes **use such a model** in real life?
- What ethical issues could arise if student marks were predicted by a machine learning model?

Student Exam Performance

A1) Identify Key Elements

- **Features (inputs):** Hours studied (X1), Attendance% (X2)
- **Label (output):** Marks (Y)
- **Task type: Regression** (Y is a continuous number)

A2) Train/Test Split

- With 6 rows and an **80/20 split** → **5 rows train, 1 row test**.
- **Why split?** To estimate generalization on **unseen** data and avoid overfitting.

A3) Algorithm Choice

- Reasonable: **Linear Regression, Decision Tree Regressor, SVR (regression)**.
- **Best first choice:** Linear Regression (relationship appears roughly linear, features are numeric, small dataset).

A4) Example Model & Predictions

- Using a simple multiple linear regression on the 6 points (intercept + X1 + X2), we get an approximate fitted model:

A5) Evaluation Metrics

- **Use: MSE, MAE, R²** (regression metrics).
- **Not appropriate:** Accuracy/Precision/Recall (classification metrics).

A6) Likely Challenges

- **Small data** → unstable coefficients, high variance.
- **Collinearity** between Hours & Attendance.
- **Omitted variables** (prior knowledge, sleep, health, quality of study).
- **Potential overfitting** with complex models on so few rows.


A7) Real-World Uses & Ethics

- **Uses:** Early warnings for at-risk students, personalized study plans, counseling resource allocation.
- **Ethics:** Avoid labeling students deterministically; ensure transparency, opt-out, and human oversight; prevent bias against students with constraints (work, health)

Attendance gets a **slightly negative** coefficient here (counter-intuitive) because the dataset is **tiny** and **Hours and Attendance are highly correlated** (multicollinearity). With **more data** or regularization, we expect Attendance to have a **positive** effect. This is a great example of why we validate assumptions and collect more data.


Unsupervised Learning

- Unsupervised learning is a branch of **ML** that deals with unlabeled data.
- Unsupervised learning algorithms **are tasked with finding patterns and relationships within the data without any prior knowledge of the data's meaning.**
- These algorithms **find hidden patterns and data without any human intervention, i.e., we don't give output to our model.**
- The training model has **only input parameter values and discovers the groups or patterns on its own.**



CLUSTERING

The process of grouping data points into clusters based on their similarity.



Example: Segmenting customers into groups based on purchase behavior

ASSOCIATION RULE LEARNING


A technique for discovering relationships between items in a dataset.



Example: "If a customer buys bread, they are also likely to buy butter."

DIMENSIONALITY REDUCTION

The process of reducing the number of features in a dataset while preserving as much information as possible.



Example: Reducing a face image with thousands of pixel features to a smaller set of key features

The image shows set of animals: elephants, camels, and cows that represents raw data that the unsupervised learning algorithm will process.

- The **"Interpretation"** stage signifies that the algorithm doesn't have predefined labels or categories for the data. It needs to figure out how to group or organize the data based on inherent patterns.
- **Algorithm** represents the core of unsupervised learning process using techniques like clustering, dimensionality reduction, or anomaly detection to identify patterns and structures in the data.
- **Processing** stage shows the algorithm working on the data.

Working of Unsupervised Learning


- Unsupervised learning works by **analyzing unlabeled data** to identify patterns and relationships.
 - Unstructured data:** May contain noisy(meaningless) data, missing values, or unknown data
 - Unlabeled data:** Data only contains a value for input parameters, there is no targeted value(output).
- The data is not labeled with any predefined categories or outcomes, so the algorithm must find these patterns and relationships on its own.

1. Data-set in Figure A is Mall data that contains information about its clients that subscribe to them.
2. Once subscribed they are provided a membership card and the mall has complete information about the customer and his/her every purchase.
3. Now using this data and unsupervised learning techniques, the mall can easily group clients based on the parameters we are feeding in.

CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)
1	Male	19	15	39
2	Male	21	15	81
3	Female	20	16	6
4	Female	23	16	77
5	Female	31	17	40
6	Female	22	17	76
7	Female	35	18	6
8	Female	23	18	94
9	Male	64	19	3
10	Female	30	19	72
11	Male	67	19	14
12	Female	35	19	99
13	Female	58	20	15
14	Female	24	20	77
15	Male	37	20	13
16	Male	22	20	79
17	Female	35	21	35

Figure A

Types of Unsupervised Learning



Clustering: It is the process of grouping data points into clusters based on their similarity.

This technique is useful for identifying patterns and relationships in data without the need for labeled examples.

In a shopping mall, customers' purchase behaviors (amount spent, frequency, type of products) can be clustered.

- Cluster 1: Students buying low-cost items.
- Cluster 2: Families buying groceries weekly.
- Cluster 3: Professionals buying luxury products.

This helps in **customer segmentation** for marketing.

Association rule learning: It is a technique for discovering relationships between items in a dataset.

It identifies rules that indicate the presence of one item implies the presence of another item with a specific probability.

In a supermarket dataset:

- Rule: If a customer buys **bread**, they are also likely to buy **butter** with 70% probability.
- Rule: If a customer buys **diapers**, they often buy **baby wipes**.

This is the basis of **Market Basket Analysis** used in recommendation systems ("Customers who bought this also bought...").

Dimensionality Reduction: Dimensionality reduction is the process of reducing the number of features in a dataset while preserving as much information as possible.

This technique is useful for improving the performance of machine learning algorithms and for data visualization.

• A facial recognition system may capture **thousands of pixel features** from an image. Instead of using all, **Principal Component Analysis (PCA)** reduces it to a smaller set of important features (like edges, shapes, key points).

• Another case: In genetics, reducing **10,000+ gene expression features** to a handful of principal components for easier analysis.

This helps in **speeding up machine learning models** and **visualizing high-dimensional data** in 2D/3D plots.

Advantages & Disadvantages of Unsupervised Learning

Challenges of Unsupervised Learning

1. **Noisy Data:** Outliers and noise can distort patterns and reduce the effectiveness of algorithms.
2. **Assumption Dependence:** Algorithms often rely on assumptions (e.g., cluster shapes), which may not match the actual data structure.
3. **Overfitting Risk:** Overfitting can occur when models capture noise instead of meaningful patterns in the data.
4. **Limited Guidance:** The absence of labels restricts the ability to guide the algorithm toward specific outcomes.
5. **Cluster Interpretability:** Results, such as clusters, may lack clear meaning or alignment with real-world categories.
6. **Sensitivity to Parameters:** Many algorithms require careful tuning of hyperparameters, such as the number of clusters in k-means.
7. **Lack of Ground Truth:** Unsupervised learning lacks labeled data, making it difficult to evaluate the accuracy of results.

Advantages of Unsupervised Machine Learning

- It helps to discover hidden patterns and various relationships between the data.
- Used for tasks such as **customer segmentation**, **anomaly detection**, and **data exploration**.
- It does not require labeled data and reduces the effort of data labeling.


Disadvantages of Unsupervised Machine Learning

- Without using labels, it may be difficult to predict the quality of the model's output.
- Cluster Interpretability may not be clear and may not have meaningful interpretations.
- It has techniques such as autoencoders and dimensionality reduction that can be used to extract meaningful features from raw data.

Reinforcement Learning

- Reinforcement learning is a **feedback-based learning method**, in which a learning agent gets a **reward for each right action** and gets a **penalty for each wrong action**.
- The agent learns automatically with these feedbacks and improves its performance.
- This algorithm is a learning method that **interacts** with the **environment** by producing actions and discovering errors.
- **Trial, error, and delay** are the most relevant characteristics of reinforcement learning.
- In this technique, the model keeps on increasing its performance using **Reward Feedback** to **learn the behavior or pattern**.

Important Note: This learning is similar to supervised as the correct o/p values are known for each I/P pattern but in some cases less information is available, eg. A network might told that its actual O/P is only 50% correct i.e. only critic information is present, hence learning based on this critic information is called **Reinforcement learning**.



Reinforcement Learning

- Unlike supervised learning, which learns from pre-collected data to make predictions or inferences, RL focuses on building an interaction between a model (agent) and its environment and data.
- It's not about training on a static dataset but about a dynamic environment where the model interacts to learn.
- The agent takes actions in the environment, and the environment, in turn, provides observations (new states) and reward signals back to the agent, creating a continuous feedback loop.

Learning Goal: Maximizing Desirable States/Cumulative Rewards:


- The primary goal in RL is to **learn how to take actions in an environment such that desirable states are achieved**, which are measured through **rewards** given to the models over time.
- The agent aims to **accumulate a total reward throughout its entire "lifetime"** (from time one to capital T).
- This often involves a **discounted sum of future rewards**, meaning current rewards are typically valued more highly than future ones, encouraging the model to prioritize sooner rewards

Data Structure and Learning Mechanism:


RL moves away from the "X's and Y's" or just "X's" paradigm seen in supervised and unsupervised learning. Instead, it deals with "**state-action pairs**".

- **States** are the observations the model sees while interacting with the environment.
- **Actions** are how the model interacts with the environment and makes decisions to receive new observations.

Feedback and Optimization: The environment provides a **reward signal** that indicates the desirability of the state the agent has entered. This signal is crucial for learning and is used to update the agent's policy or value function.



Types of Reinforcement Learning



- Rewards the agent for taking a desired action.
 - Encourages the agent to repeat the behavior.
- Eg.:** Giving a treat to a dog for sitting, providing a point in a game for a correct answer.

- Removes an undesirable stimulus to encourage a desired behavior.
 - Discourages the agent from repeating the behavior.
- Eg. :** Turning off a loud buzzer when a lever is pressed, avoiding a penalty by completing a task.

Advantages/Disadvantages

- It has autonomous decision-making that is well-suited for tasks and that can learn to make a sequence of decisions, like robotics and game-playing.
- This technique is preferred to achieve long-term results that are very difficult to achieve.
- It is used to solve a complex problems that cannot be solved by conventional techniques.

- Training Reinforcement Learning agents can be computationally expensive and time-consuming.
- Reinforcement learning is not preferable to solving simple problems.
- It needs a lot of data and a lot of computation, which makes it impractical and costly.

Applications

- **Game Playing:** RL can teach agents to play games, even complex ones.
- **Robotics:** RL can teach robots to perform tasks autonomously.
- **Autonomous Vehicles:** RL can help self-driving cars navigate and make decisions.
- **Recommendation Systems:** RL can enhance recommendation algorithms by learning user preferences.
- **Healthcare:** RL can be used to optimize treatment plans and drug discovery.
- **Finance and Trading:** RL can be used for algorithmic trading.
- **Supply Chain and Inventory Management:** RL can be used to optimize supply chain operations.
- **Energy Management:** RL can be used to optimize energy consumption.
- **Game AI:** RL can be used to create more intelligent and adaptive NPCs in video games.
- **Adaptive Personal Assistants:** RL can be used to improve personal assistants.
- **Virtual Reality (VR) and Augmented Reality (AR):** RL can be used to create immersive and interactive experiences.
- **Industrial Control:** RL can be used to optimize industrial processes.